

Can RL Improve Generalization of LLM Agents?

An Empirical Study

Zhiheng Xi^{1*}, Xin Guo^{1*}, Jiaqi Liu¹, Jiazheng Zhang¹, Yutao Fan³, Zhihao Zhang¹,
Shichun Liu¹, Mingxu Chai¹, Xiaowei Shi², Yitao Zhai², Xunliang Cai², Tao Gui^{1†},
Qi Zhang^{1†}, Xuanjing Huang^{1†}

¹Fudan University ²Meituan ³Shanghai Artificial Intelligence Laboratory
zhxi22@m.fudan.edu.cn, {tgui,qz,xjhuang}@fudan.edu.cn

Reinforcement fine-tuning (RFT) has shown promise for training LLM agents to perform multi-turn decision-making based on environment feedback. However, most existing evaluations remain largely in-domain—training and testing are conducted in the same environment or even on the same tasks. In real-world deployment, agents may operate in unseen environments with different background knowledge, observation spaces, and action interfaces. To characterize the generalization profile of RFT under such shifts, we conduct a systematic study along three axes: (1) within-environment generalization across task difficulty, (2) cross-environment transfer to unseen environments, and (3) sequential multi-environment training to quantify transfer and forgetting. Our results show that RFT generalizes well across task difficulty within an environment, but exhibits weaker transfer to unseen environments, which correlates with shifts in both semantic priors and observation/action interfaces. In contrast, sequential training yields promising downstream gains with minimal upstream forgetting, and mixture training across environments improves the overall balance. We further provide detailed analyses and deeper insights, and hope our work helps the community develop and deploy generalizable LLM agents.

1. Introduction

Reinforcement fine-tuning (RFT) has emerged as a promising post-training paradigm for improving large language model (LLM) agents on complex interactive tasks such as web navigation and software engineering (Deng et al., 2023; He et al., 2024; Jimenez et al., 2023; Merrill et al., 2026; Zan et al., 2025; Zhou et al., 2023). In RFT, an agent is trained to make a sequence of intelligent decisions that maximize task-specific objectives based on environment feedback (Mai et al., 2025; Xi et al., 2025c). By optimizing for long-horizon outcomes, RFT is often associated with stronger agentic behaviors such as instruction following (Bai et al., 2022b; Ouyang et al., 2022a), planning (Shinn et al., 2023; Yao et al., 2022b), reasoning (Lightman et al., 2023; Uesato et al., 2022), and tool use (Li et al., 2023, 2025b).

Despite rapid progress, most empirical evidence for RFT focuses on in-domain evaluation, where training and testing are conducted within the same environment—and often even on closely overlapping tasks (Liu et al., 2023; Zhou et al., 2023). In real-world deployment, however, agents frequently encounter unseen environments that differ not only in task instances, but also in background knowledge, observation spaces, and action interfaces (Ruan et al., 2023; Xi et al., 2025a). Such shifts can fundamentally reshape the interaction dynamics faced by agents, including which observations are informative, what actions are feasible, and how failures can be corrected. This raises a practical and

*Equal contribution. †Corresponding authors.

important research question: do the improvements brought by RFT generalize beyond the training distribution?

To bridge this gap, we conduct a systematic study to explore how RFT impacts the generalization and transferability of LLM agents across three axes (Figure 1). First, we investigate task-level generalization under the same environment by reviewing varying difficulty performance when RFT on a specific subset of tasks (Section 4); Second, we examine environment-level generalization to assess whether RFT maintains efficacy in unseen environments, where agents must navigate shifts in background knowledge alongside changes in observation and action spaces (Section 5). Third, we analyze sequential training across environments to characterize the dynamics of transfer and forgetting, comparing this approach against joint training on environmental mixtures (Section 6). Together, these three axes provide a comprehensive, systematic framework for understanding RFT-driven generalization and transfer in LLM agents.

Our analysis covers five agentic environments that vary in background and core properties, as shown in Table 1 and Table 4. We highlight several key findings: (1) Intra-environment generalization, within a consistent environment, RFT-trained agents exhibit significant generalization, and easy-to-hard curriculum learning further boosts performance gains. (2) Inter-environment sensitivity, while RFT enhances agentic capabilities, its generality to unseen environments exhibits fluctuations. Our analysis reveals that this across-environment generalization is sensitive to the required prior knowledge, observation spaces, and action spaces. (3) Multi-environment training. Employing multi-environment training strategies substantially enhances generalization: sequential RFT enables effective transfer to new environments while maintaining performance on upstream ones, and mix RFT across multiple environments achieves consistently strong results overall. We further provide insights that shed light on the mechanisms underlying these behaviors. Overall, we hope these findings inform the development of more generalizable LLM agents and their practical deployment in real-world settings.




Figure 1 | An overview of three axes we study.

Table 1 | Characteristics of each environment we study. The columns “DEN.”, “VAL.”, “ACT.”, “KNWL.” and “STR.” indicate whether the rewards are *dense*, action validation is *strict*, *valid action lists* are provided per step, *world knowledge* is required, and observations are *structured*, respectively. The environments vary along these dimensions.

Environment	Types	DEN.	VAL.	ACT.	KNWL.	STR.
WebShop	Web	✓	×	×	✓	✓
SearchQA	Search	×	✓	×	✓	✓
TextCraft	Game	×	✓	×	×	×
Alf World	Household	×	✓	×	×	×
BabyAI	Embodied	✓	×	✓	×	×

2. Related Work

2.1. RL for LLM Agents

Reinforcement learning (RL) has been widely adopted in LLM training, with paradigms such as reinforcement learning from human feedback (RLHF) (Bai et al., 2022a; Christiano et al., 2017; Ouyang et al., 2022b) and reinforcement learning from verifiable rewards (RLVR) (DeepSeek-AI, 2025; Rafailov et al., 2023; Schulman et al., 2017) demonstrating strong effectiveness in improving instruction following, reasoning accuracy, and behavioral alignment (Guo et al., 2025; Mu et al., 2024; Ren et al., 2025). Building upon these advances, recent work has extended reinforcement learning to LLM agents to enhance multi-step decision making (Wang et al., 2025b; Zhai et al., 2024, 2025), long-horizon planning (Song et al., 2024a; Wang et al., 2025a; Xi et al., 2025c), and interaction with external environments (Chen et al., 2025b; Feng et al., 2024; Tan et al., 2024). Across diverse agent settings, RL has been shown to strengthen agents’ abilities to decompose complex tasks, coordinate reasoning with action, and adapt policies through environment feedback (Tian et al., 2024), enabling more effective information gathering (Jin et al., 2025; Ramrakhya et al., 2025; Zhang et al., 2025b), iterative self-correction (Kumar et al., 2025; Ma et al., 2025; Zeng et al., 2025b), and tool utilization (Feng et al., 2025; Li et al., 2025a; Zeng et al., 2025a). Despite these successes, most RL-based agent methods are evaluated primarily under in-domain settings with similar task distributions and interfaces, leaving open the question of how well such learned decision-making policies generalize across tasks or transfer to unseen environments.

2.2. Generalization and Forgetting by Post-training

Recent work has examined how post-training strategies affect the generalization and forgetting behaviors of LLMs (Tu et al., 2025; Zhang et al., 2025c; Zhao et al., 2024). Supervised fine-tuning (SFT) is effective for improving in-distribution performance but often leads to over-specialization and degradation of general capabilities due to representation and distributional drift (Kotha et al., 2024; Kumar et al., 2022; Luo et al., 2023; Wang et al., 2024). In contrast, RL, particularly RLVR, tends to better preserve pre-trained representations by optimizing trajectory-level objectives rather than directly fitting target distributions, enabling more robust transfer across tasks and domains (Chen et al., 2025a; Cheng et al., 2025; Chu et al., 2025; Huan et al., 2025). However, RL can also induce negative interference and winner-take-all dynamics, reducing behavioral coverage and leading to reasoning boundary shrinkage (Hu et al., 2025; Nguyen et al., 2025; Sun et al., 2025). Notably, most existing studies focus on static and single-turn LLM tasks, whereas we extend the study of post-training generalization and forgetting to multi-turn LLM agents and systematically evaluate these effects across different environments with distinct observation and action spaces.

3. Preliminaries

3.1. Preliminaries

Task Formulation. We follow the ReAct interaction paradigm (Yao et al., 2022b) to formulate a multi-turn decision-making task. The task can be represented by the tuple $(\mathcal{U}, \mathcal{S}, \mathcal{A}, \mathcal{O}, \mathcal{T}, \mathcal{R})$. Here, $\mathcal{U}$ denotes the instruction space, $\mathcal{S}$ is the state space, $\mathcal{A}$ is the action space, and $\mathcal{O}$ represents the observation space. The function $\mathcal{T} : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ denotes the deterministic state transition function, and $\mathcal{R} : \mathcal{U} \times \mathcal{S} \rightarrow \mathbb{R}$ is the reward function.

Given a task instruction $u \in \mathcal{U}$, the agent operates based on a policy π_θ (an LLM parameterized by θ). In each interaction **step t** , the state s_t consists all previous dialogues and their resulting

observations, i.e., $s_t = (a_0, o_0, \dots, a_{t-1}, o_{t-1})$. The agent generates an action $a_t \sim \pi_\theta(\cdot | u, s_t)$. The action a_t includes an internal reasoning trace and an interaction to the environment. The agent then receives an observation $o_k \in \mathcal{O}$ from the environment. This interaction loop continues until the task is completed or the maximum number of turns is reached, resulting in a complete trajectory:

$$\tau = \{u, (a_0, o_0), (a_1, o_1), \dots, (a_T, o_T)\} \quad (1)$$

The environment provides a terminal reward $\mathcal{R}(\tau) \in [0, 1]$ upon task completion.

Policy Gradient. In RL, our objective is to optimize the policy parameters θ to maximize the expected cumulative reward over all possible trajectories for the given task (Sutton and Barto, 1998):

$$\max_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [\mathcal{R}(\tau)] \quad (2)$$

To optimize the objective function $J(\theta)$, we utilize policy gradient methods (Sutton et al., 1999). Unlike value-based methods that approximate a value function, policy gradient methods directly search the policy parameter space. The core idea is to perform gradient ascent to update θ , increasing the probability of trajectories that yield high rewards. The vanilla policy gradient is formulated as:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[\mathcal{R}(\tau) \sum_{t=0}^T \nabla_{\theta} \log \pi_\theta(a_t | s_t) \right] \quad (3)$$

Based on this gradient estimation, the parameters are updated via $\theta_{new} = \theta_{old} + \alpha \nabla_{\theta} J(\theta)$, where α is the learning rate. However, standard policy gradient methods often suffer from high variance, leading to unstable training dynamics. To address this challenge, we adopt the widely-used GRPO algorithm (DeepSeek-AI, 2025), detailed in Appendix C.

3.2. Experimental Settings

Environment Setup. We select five representative agent environments, including WebShop (Yao et al., 2022a), SearchQA (Dunn et al., 2017), TextCraft (Sanghi et al., 2022), AlfWorld (Shridhar et al., 2020), and BabyAI (Chevalier-Boisvert et al., 2018), with characteristics of each environment summarized in Table 1. Among them, WebShop is an interactive web shopping environment; SearchQA is a Q&A environment augmented with context from a search engine; ALFWorld requires agents to explore rooms and execute tasks in a household setting; BabyAI is an interactive grid world simulator; and TextCraft is a game environment for crafting items in Minecraft. The task data $\mathcal{U}$ is sourced from AgentGym (Xi et al., 2025b), with detailed statistics and action spaces for each environment provided in Appendix B.

Training Setup. We perform RFT training using the AgentGym-RL framework¹ (Xi et al., 2025c) with Qwen2.5-3B-Instruct and Qwen2.5-7B-Instruct (Team, 2024). Following the ReAct (Yao et al., 2022b) paradigm, each action receives real-time environment feedback. For all experiments, we sample 8 trajectories per task and set the maximum response length to 8192 tokens. Following Xi et al. (2025b), we set the maximum interaction turns K for each environment as follows: 10 for WebShop, 5 for SearchQA, 30 for AlfWorld, 10 for BabyAI, and 15 for TextCraft.

¹<https://github.com/woooodyy/AgentGym-RL>

Table 2 | Results of generalization across task difficulties within the same environment.

Models	WebShop			SearchQA			TextCraft			Alf World			BabyAI		
	easy	hard	all	easy	hard	all	easy	hard	all	easy	hard	all	easy	hard	all
Qwen2.5-3B-Instruct															
base model	21.7	10.6	15.3	63.7	6.5	23.7	23.7	2.3	14.5	26.8	8.0	13.2	71.5	48.0	61.6
train w/ $\mathcal{U}_{\text{easy}}$	90.3	75.3	81.6	82.7	16.9	36.6	97.6	49.4	76.9	93.2	82.0	85.1	93.4	79.0	87.3
train w/ $\mathcal{U}_{\text{hard}}$	86.1	84.5	85.2	72.5	19.4	35.3	95.0	42.2	72.3	96.1	92.9	93.8	92.3	77.3	86.0
train w/ $\mathcal{U}$	92.8	84.3	87.9	87.6	22.1	41.8	93.9	47.4	73.9	97.0	89.8	91.8	93.2	77.5	86.6
$\mathcal{U}_{\text{easy}} + \mathcal{U}_{\text{hard}}$	93.6	85.2	88.7	82.1	19.8	38.5	93.9	52.9	76.3	97.3	93.4	94.4	94.5	82.2	89.3
$\mathcal{U}_{\text{hard}} + \mathcal{U}_{\text{easy}}$	90.2	82.4	85.7	82.0	18.7	37.7	97.8	40.1	73.0	95.5	92.3	93.6	94.5	80.4	88.6
Qwen2.5-7B-Instruct															
base model	44.1	17.4	28.6	79.6	10.4	31.2	46.9	16.0	33.6	40.2	21.4	26.6	80.1	47.9	67.0
train w/ $\mathcal{U}_{\text{easy}}$	88.1	77.5	82.0	85.8	21.1	40.5	98.2	47.4	76.4	95.2	89.7	91.2	95.2	76.1	87.2
train w/ $\mathcal{U}_{\text{hard}}$	87.9	81.4	84.2	93.3	27.0	46.9	99.3	57.3	81.3	97.5	94.6	95.4	92.8	77.9	86.5
train w/ $\mathcal{U}$	92.9	81.9	86.5	91.8	26.6	46.1	95.2	52.3	80.9	97.7	97.2	92.0	95.6	74.5	88.8
$\mathcal{U}_{\text{easy}} + \mathcal{U}_{\text{hard}}$	90.6	77.5	83.0	89.8	25.5	44.8	98.7	64.0	83.8	96.1	92.6	93.6	95.5	82.9	90.2
$\mathcal{U}_{\text{hard}} + \mathcal{U}_{\text{easy}}$	89.1	79.1	83.3	87.6	21.5	42.3	99.8	57.0	81.4	97.7	88.5	91.1	92.5	78.2	86.5

Evaluation Setup. As summarized in Table 1, WebShop and BabyAI provide dense rewards, while SearchQA, Alf World, and TextCraft use binary rewards. We follow DeepSeek-AI (2025) and adopt exact matching, counting an action as correct only if it exactly matches the reference. To reduce randomness, we report avg@8 as the main metric, and also measure the average interaction turns ($\bar{k}$) and generated tokens ($\bar{l}$) for efficiency. For fair comparison, all agents are evaluated with the maximum number of interaction turns set to $K = 20$ during testing.

4. Does RFT Yield Generalization Across Task Difficulties within the Same Environment?

Reinforcement fine-tuning trains LLM agents through interaction with the environment, allowing them to learn its dynamics and adapt over time (Cui et al., 2025; Song et al., 2024b). Yet even with a fixed action and observation space, tasks within the same environment can differ substantially in exploration depth and information accessibility. In this section, we investigate whether RFT policies learned on a subset of tasks transfer to other tasks of differing difficulty within the same environment.

Setting. Following practice in previous work (Bengio et al., 2009; Mukherjee et al., 2023), we categorize the tasks $\mathcal{U}$ into easy and hard difficulty levels (denoted as $\mathcal{U}_{\text{easy}}$ and $\mathcal{U}_{\text{hard}}$) based on the avg@8 results of the Qwen2.5-7B-Instruct model, while ensuring a balanced distribution of data between the two difficulty levels. The same categorization is applied to the test set. Detailed data statistics for each environment are provided in Appendix B.

RFT demonstrates strong transferability across varying difficulty levels within the same environment. As shown in Table 2, RFT demonstrates strong robustness when trained on data of varying difficulty levels within the same environment. Notably, with 7B model on WebShop, training on $\mathcal{U}_{\text{easy}}$ improves performance on the hard testset by 60.1 points, suggesting that RFT encourages agents to adapt to the environment, thereby enabling performance transfer across tasks with different difficulty levels and varying steps (Chen et al., 2025a; Huan et al., 2025). Overall, we also find that training on $\mathcal{U}_{\text{hard}}$ yields more gains on test set, e.g., for 3B model in Alf World, training on $\mathcal{U}_{\text{hard}}$ outperforms training on $\mathcal{U}_{\text{easy}}$ by 8.7 points. We attribute this trend to the richer failure signals and

longer-horizon exploration induced by harder tasks, which push the RFT optimizer to seek better policies (Wang et al., 2023; Xu et al., 2023).

Curriculum learning can further enhance performance. Additionally, we investigate the impact of the training sequence using data of varying difficulty levels. Results in Table 2 indicate that training on $\mathcal{U}$ (i.e., mixture of $\mathcal{U}_{\text{easy}}$ and $\mathcal{U}_{\text{hard}}$) generally does not yield optimal performance. In most cases, training on $\mathcal{U}_{\text{easy}}$ first, followed by $\mathcal{U}_{\text{hard}}$, achieves the best results. Compared with using a single difficulty level, this easy-to-hard curriculum enables further performance improvement. For example, on BabyAI, training on $\mathcal{U}_{\text{easy}} + \mathcal{U}_{\text{hard}}$ outperforms training on $\mathcal{U}_{\text{easy}}$ and on $\mathcal{U}_{\text{hard}}$ by 2.0 and 3.3 points, respectively. This result validates the rationale and effectiveness of curriculum learning (Qi et al., 2024; Zhang et al., 2025a).

RFT training encourages agents to perform efficient exploration.

Figure 2 reports the average number of interaction turns and generated tokens across different environments, with a more detailed breakdown provided in Figure 8 of Appendix E. We find that RFT enables agents to interact and explore more efficiently within the same environment. For example, when training the 3B model on $\mathcal{U}_{\text{easy}}$ in BabyAI, the average interaction turns decrease from 10.76 to 4.19, and the average trajectory length is reduced from 624.58 to 160.60 tokens. These results indicate that RFT not only improves success rates but also encourages more concise and goal-directed exploration, substantially enhancing interaction efficiency (Nakano et al., 2021; Song et al., 2024b). A case can be found in Appendix G.




Figure 2 | Average generated tokens and interactive turn of all environments for Qwen2.5-3B-Instruct model trained with varying-difficulty tasks.

5. Does RFT Yield Generalization Across Different Environments?

In real-world scenarios, agents may encounter previously unseen environments and tasks, which is essential for practical deployment. We therefore investigate *whether RFT improves performance in unseen environments with different background knowledge, observation space and action spaces*. Concretely, we perform RFT training within a single environment and then evaluate the trained agent on other environments to measure cross-environment generalization.

Setting. We employ the base model (i.e., Qwen2.5-3B-Instruct or Qwen2.5-7B-Instruct) as the baseline and evaluate its avg@8 metric across varying environments. To make these comparisons explicit, we define three metrics: $\Delta\text{Held-In}$, $\Delta\text{Held-Out}$, and $\Delta\text{Overall}$. $\Delta\text{Held-In}$ measures the average improvement over the baseline when the training and test environments coincide. $\Delta\text{Held-Out}$ measures the average improvement when the agent is evaluated on environments different from those seen during training. Finally, $\Delta\text{Overall}$ reports the average improvement over the baseline across all evaluation settings.

Performance differs significantly between held-in and held-out environments. Table 3 reveals that agents exhibit generalization across different environments, yet a significant performance gap exists between held-in and held-out settings. Specifically, substantial gains are achieved under held-

Table 3 | Results of generalization across different environments. Darker red indicates greater performance improvement, while darker blue signifies more pronounced performance decline. The best result is in **bold**, while the second-best is marked with underline.

Models	WebShop	SearchQA	TextCraft	AlfWorld	BabyAI	Δ Held-In	Δ Held-Out	Δ Overall
<i>Qwen2.5-3B-Instruct</i>								
base model	15.30	23.66	14.50	13.19	61.83	–	–	–
train w/ WebShop	87.86	25.84	18.50	23.06	70.91	+72.56	+6.29	+19.54
train w/ SearchQA	22.97	41.78	22.75	12.13	65.72	+18.13	+4.69	+ 4.56
train w/ TextCraft	14.46	22.16	73.88	11.00	63.73	+59.38	–0.66	+11.35
train w/ AlfWorld	21.86	23.50	24.88	91.81	64.68	+78.62	+4.91	+19.65
train w/ BabyAI	28.36	22.63	16.75	4.50	86.55	+24.72	+1.40	+ 6.06
<i>Qwen2.5-7B-Instruct</i>								
base model	28.59	31.19	33.63	26.56	67.00	–	–	–
train w/ WebShop	86.50	33.28	40.75	24.13	79.21	+57.91	+4.75	+15.38
train w/ SearchQA	47.07	46.12	35.25	16.75	80.33	+14.93	+5.91	+ 7.71
train w/ TextCraft	38.30	32.19	80.88	31.50	77.95	+47.25	+6.65	+14.77
train w/ AlfWorld	34.31	29.59	36.13	92.00	72.91	+65.44	+3.13	+15.59
train w/ BabyAI	10.25	29.41	39.25	28.13	88.79	+21.79	–3.23	+ 1.77

in conditions. In AlfWorld, performance improves by 78.62 and 65.44 points for the 3B and 7B models, respectively. In contrast, in held-out conditions, generalization remains possible in most cases, but with more modest gains. On average, the 3B and 7B models yield improvements of 3.32 and 3.44 points on unseen environments, respectively.

In unseen environments, positive transfer is observed in most cases. Compared to the baseline, most trained agents can demonstrate transferability to unseen environments, even when these environments demand different background knowledge and operate under different action spaces. Notably, agents trained on WebShop, AlfWorld, and SearchQA consistently exhibit positive Δ held-out. When evaluated on WebShop, agents trained on SearchQA yield performance gains of 7.67 and 18.48 points for the 3B and 7B models, respectively. This can be attributed to the similarity between WebShop and SearchQA as search-based environments. Specifically, the agent trained on SearchQA learns to formulate more flexible search queries, as well as efficient information extraction from complex results. An illustrative case is presented in Figure 10 in Appendix G and discussed in Section 7.

In unseen environments, negative effects may sometimes occur. However, agents trained on TextCraft and BabyAI may struggle to generalize to other environments, e.g., after training on BabyAI, the 7B models show average negative improvements of –3.23 on held-out tasks, and even drop sharply from 28.59 to 10.25 on WebShop. Through careful analysis, we find that because the BabyAI environment provides available actions at each step, the agent gradually becomes dependent on this information during training, leading to a decline in long-horizon reasoning capability. When faced with other environments, it fails to accurately use the valid action, resulting in a sharp performance drop. A case of interaction between the agent and BabyAI environment is provided in Appendix G.

Cross-environment generalization performance varies significantly across target environments. Moreover, our observations reveal that cross-environment generalization performance is highly dependent on the target environment. On the one hand, agents trained on nearly all source




Figure 3 | Training dynamics of forgetting and transfer in sequential two-stage cross-environment training with Qwen2.5-7B-Instruct, where blue and red denote the upstream environment and the downstream environment, respectively.

environments generalize effectively to TextCraft and BabyAI, as these domains rely less on specific background knowledge, thereby allowing acquired skills to transfer more readily. On the other hand, effective transfer to Alf World and SearchQA proves to be significantly more challenging. We attribute this to the strict action validation and sparse feedback inherent in these two environments. For instance, Alf World responds uniformly with “Nothing happens.” to all invalid actions, offering no instructional guidance for improvement.

6. How Does Sequential RFT Across Environments Affect Transfer and Forgetting?

Beyond zero-shot transfer to unseen environments without additional training, we further investigate the effect of sequential training across multiple environments on agentic behaviors, specifically the resulting learning dynamics on memorizing downstream tasks or forgetting on upstream ones. In our experimental setup, a model initially converged in one environment is further trained on a second environment, after which we assess both its retention of performance on the original environment and its adaptation to the new one. Finally, we extend this analysis to sequential training across five environments, comparing this strategy with joint training across a mixture of environments.

Setting. We conduct 20 two-stage experiments by sequentially pairing each of the five environments as the upstream and the downstream, as well as several five-stage training experiments. Using agents trained on single-environment as baselines, we evaluate anti-forgetting by upstream performance retention and transferability by downstream performance gains. Furthermore, we compare the five-stage sequential approach with joint training, in which the agent is trained on all available data from every environment in a randomly mixed manner.

Sequential training demonstrates consistent anti-forgetting and transferability. Figure 3 illustrates the performance dynamics of 8 two-stage sequential training scenarios, with others and the final results summarized in Figure 7 and Table 6 in Appendix D. Overall, sequential training matches or exceeds single-task performance on the downstream environment while largely preserving upstream performance. For instance, when the WebShop-pre-trained agent is further trained on TextCraft, it boosts performance on TextCraft from the single-task baseline of 80.88 to 82.50. Meanwhile, performance on WebShop experiences only a minor fluctuation, shifting from 86.5 to 86.32.




Figure 4 | Generalization of sequential two-stage cross-environment training with Qwen2.5-7B-Instruct. Across five environments (WS, SQ, TC, AW, BA), the figure presents the generalization performance on the *three unseen environments* following sequential training on two environments. Each subplot corresponds to a fixed first environment, with *dashed line* indicating the baseline performance.

These findings suggest that sequential training endows agents with stable capabilities for transfer and resistance to forgetting.

Generalization in multi-environment training is highly correlated with single-environment training. Figure 4 demonstrates the generalization performance of multi-environment training, revealing a strong alignment with the generalization patterns in single-environment training. First, environments that yield poor generalization in single-environment settings continue to be detrimental. For instance, the agent trained on BabyAI exhibits severe negative effect on WebShop. Even when BabyAI is trained sequentially after other generalizable environments (i.e., TC+BA, AW+BA), it drastically degrades their WebShop performance. Notably, the score of TC+BA dropped sharply from 38.30 (score of TC in Table 3) to 23.0.

Moreover, from the perspective of target environments, the results also show consistency with single-environment generalization. For target environments easy to generalization in single-environment scenarios, such as TextCraft and BabyAI, sequentially trained agents also tend to perform well. Specifically, agents pre-trained on WebShop (i.e., WS+SQ, WS+AW, WS+BA) achieve significant gains on TextCraft, increasing by 11.00, 10.12, and 15.62 points, respectively. However, for environments challenging for generalization, like Alf World, sequential training yields limited benefits. For instance, all three agents utilizing TextCraft as the upstream environment (TC+WS, TC+SQ, TC+BA) suffer performance degradation on Alf World.

Training order significantly affects generalization performance in held-out environments. As shown in Figure 4, the training order exerts a substantial influence on generalization performance. For instance, when evaluated on TextCraft and Alf World, BA+SQ outperforms SQ+BA by 15.63 and 10.00 points, respectively. This represents a substantial margin, particularly for held-out scenarios. We attribute this phenomenon to the inherent task difficulty within each environment. As noted in Section 5, BabyAI provides detailed feedback, whereas SearchQA imposes strict constraints with




Figure 5 | Training dynamics of sequential training across five environments. We present the results for representative sequence combinations, monitoring how performance on each environment changes as the agent is trained on different environments sequentially. The *dashed lines* denote the performance achieved by **joint training on a mixture of data from all five environments**.

limited feedback. Consequently, the BA+SQ order naturally creates an “easy-to-hard” curriculum, which in turn facilitates better generalization performance.

Sequential training achieves performance comparable to joint training. Furthermore, we conduct sequential training over five environments in different orders. Figure 5 illustrates the performance dynamics on each environment across the five training stages for two representative sequences; for comparison, joint-training results are indicated by dashed lines. The results show that sequential training achieves performance comparable to joint training, even after training on five distinct tasks. Overall, the final performance is insensitive to the training order, which we attribute to RFT’s ability to preserve previously acquired capabilities, consistent with findings in prior work. For environments such as AlfWorld and SearchQA—to which other tasks struggle to generalize—relatively pronounced forgetting may occur over the course of long-term sequential training.

7. Further Analysis and Discussion

Failure Mode Analysis. We conduct a fine-grained error analysis on both intra/inter-environment evaluations, summarizing 8 common failure modes including *instruction misinterpretation*, *action execution failure*, and *logical deficits*. Using GPT-5-mini, we systematically categorize the error trajectories across all scenarios. More detail about each failure pattern can be found in Appendix F.

The results presented in Figure 6 reveal that while failure modes differ by environment, errors related to “Confirmation Bias” are prevalent (> 10%) across all scenarios. This suggests that after training, agents tend to exhibit overconfidence, neglecting further verification and lacking the capacity for self-reflection based on environmental feedback. Moreover, in SearchQA, errors categorized as “Guessing or Fabrication” are widespread among both held-in (23.8%) and held-out (21.2%) settings. This highlights a critical deficiency in tool utilization, which fundamentally constrains the agent’s generalization potential.

When the training and test environments differ, agents exhibit different failure modes. For example, in WebShop we observe an increase in the proportion of “State or Memory Inconsistency” errors, rising from 4.3% in held-in to 21.9% in held-out settings. This suggests that, when confronted with large volumes of information, the agent’s decision-making becomes less coherent, and it struggles to




Figure 6 | Comparison of failure mode distributions between in-environment and out-of-environment evaluation.

generalize its ability to extract salient information. A full environment-by-environment breakdown is provided in Figure 9.

Case Study. We perform a comprehensive analysis of environment generalization and present specific case studies in Appendix G. Notably, Figure 10 illustrates the mechanism behind the successful transfer from SearchQA to WebShop. Specifically, the base agent tends to blindly input the entire instruction into the search bar, and struggles to accurately extract key information from the voluminous HTML content, resulting in the retrieval and selection of irrelevant items. In contrast, the agent trained on SearchQA learns to effectively search for key details and extract information, successfully selecting the correct item.

Furthermore, Figure 11 offers qualitative insight into why SearchQA generalization remains difficult, using an AlfWorld-trained agent as a case study. Although both agents fail initially, the SearchQA-trained agent iteratively refines its queries to become more specific, improving retrieval and ultimately recovering. By contrast, this query-refinement behavior does not reliably transfer: the AlfWorld-trained agent falls into a degenerate loop, repeatedly issuing near-duplicate searches, returning the same responses, and ultimately failing.

8. Conclusion

In this paper, we present a systematic study of how RFT affects the transfer and generalization of LLM agents for multi-turn decision-making. Through large-scale experiments along three complementary axes, we characterize when RFT generalizes within and across environments and identify generalization patterns. We further complement our quantitative results with failure mode analysis and qualitative case study to pinpoint where agents break down and what behaviors fail to transfer. Together, our findings offer practical guidance for training and evaluating agents under distribution shift, and we hope they inform the development of agents that generalize reliably in real-world deployments.

References

- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosiute, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemí Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. Constitutional AI: harmlessness from AI feedback. *CoRR*, abs/2212.08073, 2022a. doi: 10.48550/ARXIV.2212.08073. URL <https://doi.org/10.48550/arXiv.2212.08073>.
- Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*, 2022b.
- Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In *Proceedings of the 26th annual international conference on machine learning*, pages 41–48, 2009.
- Howard Chen, Noam Razin, Karthik Narasimhan, and Danqi Chen. Retaining by doing: The role of on-policy data in mitigating forgetting. *CoRR*, abs/2510.18874, 2025a. doi: 10.48550/ARXIV.2510.18874. URL <https://doi.org/10.48550/arXiv.2510.18874>.
- Kevin Chen, Marco F. Cusumano-Towner, Brody Huval, Aleksei Petrenko, Jackson Hamburger, Vladlen Koltun, and Philipp Krähenbühl. Reinforcement learning for long-horizon interactive LLM agents. *CoRR*, abs/2502.01600, 2025b. doi: 10.48550/ARXIV.2502.01600. URL <https://doi.org/10.48550/arXiv.2502.01600>.
- Zhoujun Cheng, Shibo Hao, Tianyang Liu, Fan Zhou, Yutao Xie, Feng Yao, Yuexin Bian, Yonghao Zhuang, Nilabjo Dey, Yuheng Zha, et al. Revisiting reinforcement learning for llm reasoning from a cross-domain perspective. *arXiv preprint arXiv:2506.14965*, 2025.
- Maxime Chevalier-Boisvert, Dzmitry Bahdanau, Salem Lahlou, Lucas Willems, Chitwan Saharia, Thien Huu Nguyen, and Yoshua Bengio. Babyai: A platform to study the sample efficiency of grounded language learning. *arXiv preprint arXiv:1810.08272*, 2018.
- Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pages 4299–4307, 2017.
- Tianzhe Chu, Yuexiang Zhai, Jihan Yang, Shengbang Tong, Saining Xie, Dale Schuurmans, Quoc V Le, Sergey Levine, and Yi Ma. Sft memorizes, rl generalizes: A comparative study of foundation model post-training. *arXiv preprint arXiv:2501.17161*, 2025.
- Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, et al. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025.

- DeepSeek-AI. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *CoRR*, abs/2501.12948, 2025. doi: 10.48550/ARXIV.2501.12948. URL <https://doi.org/10.48550/arXiv.2501.12948>.
- Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Sam Stevens, Boshi Wang, Huan Sun, and Yu Su. Mind2web: Towards a generalist agent for the web. *Advances in Neural Information Processing Systems*, 36:28091–28114, 2023.
- Matthew Dunn, Levent Sagun, Mike Higgins, V Ugur Guney, Volkan Cirik, and Kyunghyun Cho. Searchqa: A new q&a dataset augmented with context from a search engine. *arXiv preprint arXiv:1704.05179*, 2017.
- Jiazhan Feng, Shijue Huang, Xingwei Qu, Ge Zhang, Yujia Qin, Baoquan Zhong, Chengquan Jiang, Jinxin Chi, and Wanjuan Zhong. Retool: Reinforcement learning for strategic tool use in llms. *CoRR*, abs/2504.11536, 2025. doi: 10.48550/ARXIV.2504.11536. URL <https://doi.org/10.48550/arXiv.2504.11536>.
- Peiyuan Feng, Yichen He, Guanhua Huang, Yuan Lin, Hanchong Zhang, Yuchen Zhang, and Hang Li. AGILE: A novel reinforcement learning framework of LLM agents. In Amir Globersons, Lester Mackey, Danielle Belgrave, Angela Fan, Ulrich Paquet, Jakub M. Tomczak, and Cheng Zhang, editors, *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*, 2024.
- Xu Guo, Tianyi Liang, Tong Jian, Xiaogui Yang, Ling-I Wu, Chenhui Li, Zhihui Lu, Qipeng Guo, and Kai Chen. Ifdecorator: Wrapping instruction following reinforcement learning with verifiable rewards. *arXiv preprint arXiv:2508.04632*, 2025.
- Hongliang He, Wenlin Yao, Kaixin Ma, Wenhao Yu, Yong Dai, Hongming Zhang, Zhenzhong Lan, and Dong Yu. Webvoyager: Building an end-to-end web agent with large multimodal models. *arXiv preprint arXiv:2401.13919*, 2024.
- Chuxuan Hu, Yuxuan Zhu, Antony Kellermann, Caleb Biddulph, Suppakit Waiwitlikhit, Jason Benn, and Daniel Kang. Breaking barriers: Do reinforcement post training gains transfer to unseen domains? *CoRR*, abs/2506.19733, 2025. doi: 10.48550/ARXIV.2506.19733. URL <https://doi.org/10.48550/arXiv.2506.19733>.
- Jian Hu. Reinforce++: A simple and efficient approach for aligning large language models. *arXiv preprint arXiv:2501.03262*, 2025.
- Maggie Huan, Yuetai Li, Tuney Zheng, Xiaoyu Xu, Seungone Kim, Minxin Du, Radha Pooven-dran, Graham Neubig, and Xiang Yue. Does math reasoning improve general LLM capabilities? understanding transferability of LLM reasoning. *CoRR*, abs/2507.00432, 2025. doi: 10.48550/ARXIV.2507.00432. URL <https://doi.org/10.48550/arXiv.2507.00432>.
- Carlos E Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swe-bench: Can language models resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2023.
- Bowen Jin, Hansi Zeng, Zhenrui Yue, Dong Wang, Hamed Zamani, and Jiawei Han. Search-r1: Training llms to reason and leverage search engines with reinforcement learning. *CoRR*, abs/2503.09516, 2025. doi: 10.48550/ARXIV.2503.09516. URL <https://doi.org/10.48550/arXiv.2503.09516>.

- Suhas Kotha, Jacob Mitchell Springer, and Aditi Raghunathan. Understanding catastrophic forgetting in language models via implicit inference. In *The Twelfth International Conference on Learning Representations, ICLR 2024, Vienna, Austria, May 7-11, 2024*. OpenReview.net, 2024. URL <https://openreview.net/forum?id=VrHiF2hsrm>.
- Ananya Kumar, Aditi Raghunathan, Robbie Matthew Jones, Tengyu Ma, and Percy Liang. Fine-tuning can distort pretrained features and underperform out-of-distribution. In *The Tenth International Conference on Learning Representations, ICLR 2022, Virtual Event, April 25-29, 2022*. OpenReview.net, 2022. URL <https://openreview.net/forum?id=UYneFzXSJWh>.
- Aviral Kumar, Vincent Zhuang, Rishabh Agarwal, Yi Su, John D. Co-Reyes, Avi Singh, Kate Baumli, Shariq Iqbal, Colton Bishop, Rebecca Roelofs, Lei M. Zhang, Kay McKinney, Disha Shrivastava, Cosmin Paduraru, George Tucker, Doina Precup, Feryal M. P. Behbahani, and Aleksandra Faust. Training language models to self-correct via reinforcement learning. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025*. OpenReview.net, 2025. URL <https://openreview.net/forum?id=CjwERcAU7w>.
- Lei Li, Yekun Chai, Shuohuan Wang, Yu Sun, Hao Tian, Ningyu Zhang, and Hua Wu. Tool-augmented reward modeling. *arXiv preprint arXiv:2310.01045*, 2023.
- Xiaoxi Li, Wenxiang Jiao, Jiarui Jin, Guanting Dong, Jiajie Jin, Yinuo Wang, Hao Wang, Yutao Zhu, Ji-Rong Wen, Yuan Lu, and Zhicheng Dou. Deepagent: A general reasoning agent with scalable toolsets. *CoRR*, abs/2510.21618, 2025a. doi: 10.48550/ARXIV.2510.21618. URL <https://doi.org/10.48550/arXiv.2510.21618>.
- Xiaoxi Li, Wenxiang Jiao, Jiarui Jin, Guanting Dong, Jiajie Jin, Yinuo Wang, Hao Wang, Yutao Zhu, Ji-Rong Wen, Yuan Lu, et al. Deepagent: A general reasoning agent with scalable toolsets. *arXiv preprint arXiv:2510.21618*, 2025b.
- Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *The Twelfth International Conference on Learning Representations*, 2023.
- Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, et al. Agentbench: Evaluating llms as agents. *arXiv preprint arXiv:2308.03688*, 2023.
- Yun Luo, Zhen Yang, Fandong Meng, Yafu Li, Jie Zhou, and Yue Zhang. An empirical study of catastrophic forgetting in large language models during continual fine-tuning. *CoRR*, abs/2308.08747, 2023. doi: 10.48550/ARXIV.2308.08747. URL <https://doi.org/10.48550/arXiv.2308.08747>.
- Ruotian Ma, Peisong Wang, Cheng Liu, Xingyan Liu, Jiaqi Chen, Bang Zhang, Xin Zhou, Nan Du, and Jia Li. S²r: Teaching llms to self-verify and self-correct via reinforcement learning. *CoRR*, abs/2502.12853, 2025. doi: 10.48550/ARXIV.2502.12853. URL <https://doi.org/10.48550/arXiv.2502.12853>.
- Xinji Mai, Haotian Xu, Zhong-Zhi Li, Weinong Wang, Jian Hu, Yingying Zhang, Wenqiang Zhang, et al. Agent rl scaling law: Agent rl with spontaneous code execution for mathematical problem solving. *arXiv preprint arXiv:2505.07773*, 2025.
- Mike A Merrill, Alexander G Shaw, Nicholas Carlini, Boxuan Li, Harsh Raj, Ivan Bercovich, Lin Shi, Jeong Yeon Shin, Thomas Walshe, E Kelly Buchanan, et al. Terminal-bench: Benchmarking agents on hard, realistic tasks in command line interfaces. *arXiv preprint arXiv:2601.11868*, 2026.

- Tong Mu, Alec Helyar, Johannes Heidecke, Joshua Achiam, Andrea Vallone, Ian Kivlichan, Molly Lin, Alex Beutel, John Schulman, and Lilian Weng. Rule based rewards for language model safety. In Amir Globersons, Lester Mackey, Danielle Belgrave, Angela Fan, Ulrich Paquet, Jakub M. Tomczak, and Cheng Zhang, editors, *Advances in Neural Information Processing Systems 38: Annual Conference on Neural Information Processing Systems 2024, NeurIPS 2024, Vancouver, BC, Canada, December 10 - 15, 2024*, 2024.
- Subhabrata Mukherjee, Arindam Mitra, Ganesh Jawahar, Sahaj Agarwal, Hamid Palangi, and Ahmed Awadallah. Orca: Progressive learning from complex explanation traces of gpt-4. *arXiv preprint arXiv:2306.02707*, 2023.
- Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, et al. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2021.
- Phuc Minh Nguyen, Chinh D. La, Duy M. H. Nguyen, Nitesh V. Chawla, Binh T. Nguyen, and Khoa D. Doan. The reasoning boundary paradox: How reinforcement learning constrains language models. *CoRR*, abs/2510.02230, 2025. doi: 10.48550/ARXIV.2510.02230. URL <https://doi.org/10.48550/arXiv.2510.02230>.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022a.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In Sanmi Koyejo, S. Mohamed, A. Agarwal, Danielle Belgrave, K. Cho, and A. Oh, editors, *Advances in Neural Information Processing Systems 35: Annual Conference on Neural Information Processing Systems 2022, NeurIPS 2022, New Orleans, LA, USA, November 28 - December 9, 2022*, 2022b.
- Zehan Qi, Xiao Liu, Iat Long Iong, Hanyu Lai, Xueqiao Sun, Wenyi Zhao, Yu Yang, Xinyue Yang, Jiadai Sun, Shuntian Yao, et al. Webrl: Training llm web agents via self-evolving online curriculum reinforcement learning. *arXiv preprint arXiv:2411.02337*, 2024.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In Alice Oh, Tristan Naumann, Amir Globerson, Kate Saenko, Moritz Hardt, and Sergey Levine, editors, *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10 - 16, 2023*, 2023.
- Ram Ramrakhya, Matthew Chang, Xavier Puig, Ruta Desai, Zsolt Kira, and Roozbeh Mottaghi. Grounding multimodal llms to embodied agents that ask for help with reinforcement learning. *arXiv preprint arXiv:2504.00907*, 2025.
- Z. Z. Ren, Zhihong Shao, Junxiao Song, Huajian Xin, Haocheng Wang, Wanjia Zhao, Liyue Zhang, Zhe Fu, Qihao Zhu, Dejian Yang, Z. F. Wu, Zhibin Gou, Shirong Ma, Hongxuan Tang, Yuxuan Liu, Wenjun Gao, Daya Guo, and Chong Ruan. Deepseek-prover-v2: Advancing formal mathematical reasoning via reinforcement learning for subgoal decomposition. *CoRR*, abs/2504.21801, 2025. doi: 10.48550/ARXIV.2504.21801. URL <https://doi.org/10.48550/arXiv.2504.21801>.

- Jingqing Ruan, Yihong Chen, Bin Zhang, Zhiwei Xu, Tianpeng Bao, Hangyu Mao, Ziyue Li, Xingyu Zeng, Rui Zhao, et al. Tptu: Task planning and tool usage of large language model-based ai agents. In *NeurIPS 2023 Foundation Models for Decision Making Workshop*, 2023.
- Aditya Sanghi, Rao Fu, Vivian Liu, Karl D. D. Willis, Hooman Shayani, Amir Hosein Khasahmadi, Srinath Sridhar, and Daniel Ritchie. Textcraft: Zero-shot generation of high-fidelity and diverse shapes from text. *CoRR*, abs/2211.01427, 2022. doi: 10.48550/ARXIV.2211.01427. URL <https://doi.org/10.48550/arXiv.2211.01427>.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *CoRR*, abs/1707.06347, 2017. URL <http://arxiv.org/abs/1707.06347>.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36:8634–8652, 2023.
- Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. Alfworld: Aligning text and embodied environments for interactive learning. *arXiv preprint arXiv:2010.03768*, 2020.
- Yifan Song, Da Yin, Xiang Yue, Jie Huang, Sujian Li, and Bill Yuchen Lin. Trial and error: Exploration-based trajectory optimization for LLM agents. *CoRR*, abs/2403.02502, 2024a. doi: 10.48550/ARXIV.2403.02502. URL <https://doi.org/10.48550/arXiv.2403.02502>.
- Yifan Song, Da Yin, Xiang Yue, Jie Huang, Sujian Li, and Bill Yuchen Lin. Trial and error: Exploration-based trajectory optimization for llm agents. *arXiv preprint arXiv:2403.02502*, 2024b.
- Yiyu Sun, Shawn Hu, Georgia Zhou, Ken Zheng, Hannaneh Hajishirzi, Nouha Dziri, and Dawn Song. Omega: Can llms reason outside the box in math? evaluating exploratory, compositional, and transformative generalization. *arXiv preprint arXiv:2506.18880*, 2025.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement learning - an introduction*. Adaptive computation and machine learning. MIT Press, 1998. ISBN 978-0-262-19398-6. URL <http://www.incompleteideas.net/book/first/the-book.html>.
- Richard S Sutton, David McAllester, Satinder Singh, and Yishay Mansour. Policy gradient methods for reinforcement learning with function approximation. *Advances in neural information processing systems*, 12, 1999.
- Weihao Tan, Wentao Zhang, Xinrun Xu, Haochong Xia, Ziluo Ding, Boyu Li, Bohan Zhou, Junpeng Yue, Jiechuan Jiang, Yewen Li, et al. Cradle: Empowering foundation agents towards general computer control. *arXiv preprint arXiv:2403.03186*, 2024.
- Qwen Team. Qwen2.5: A party of foundation models, September 2024. URL <https://qwenlm.github.io/blog/qwen2.5/>.
- Haozhe Tian, Homayoun Hamedmoghadam, Robert Shorten, and Pietro Ferraro. Reinforcement learning with adaptive regularization for safe control of critical systems. *Advances in Neural Information Processing Systems*, 37:2528–2557, 2024.
- Chengying Tu, Xuemiao Zhang, Rongxiang Weng, Rumei Li, Chen Zhang, Yang Bai, Hongfei Yan, Jingang Wang, and Xunliang Cai. A survey on LLM mid-training. *CoRR*, abs/2510.23081, 2025. doi: 10.48550/ARXIV.2510.23081. URL <https://doi.org/10.48550/arXiv.2510.23081>.

- Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. Solving math word problems with process-and outcome-based feedback. *arXiv preprint arXiv:2211.14275*, 2022.
- Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. *arXiv preprint arXiv:2305.16291*, 2023.
- Hanlin Wang, Chak Tou Leong, Jiashuo Wang, Jian Wang, and Wenjie Li. SPA-RL: reinforcing LLM agents via stepwise progress attribution. *CoRR*, abs/2505.20732, 2025a. doi: 10.48550/ARXIV.2505.20732. URL <https://doi.org/10.48550/arXiv.2505.20732>.
- Huaijie Wang, Shibo Hao, Hanze Dong, Shenao Zhang, Yilin Bao, Ziran Yang, and Yi Wu. Offline reinforcement learning for llm multi-step reasoning. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 8881–8893, 2025b.
- Xinyi Wang, Antonis Antoniadis, Yanai Elazar, Alfonso Amayuelas, Alon Albalak, Kexun Zhang, and William Yang Wang. Generalization vs memorization: Tracing language models’ capabilities back to pretraining data. *arXiv preprint arXiv:2407.14985*, 2024.
- Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, et al. The rise and potential of large language model based agents: A survey. *Science China Information Sciences*, 68(2):121101, 2025a.
- Zhiheng Xi, Yiwen Ding, Wenxiang Chen, Boyang Hong, Honglin Guo, Junzhe Wang, Xin Guo, Dingwen Yang, Chenyang Liao, Wei He, et al. Agentgym: Evaluating and training large language model-based agents across diverse environments. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 27914–27961, 2025b.
- Zhiheng Xi, Jixuan Huang, Chenyang Liao, Baodai Huang, Honglin Guo, Jiaqi Liu, Rui Zheng, Junjie Ye, Jiazhang Zhang, Wenxiang Chen, Wei He, Yiwen Ding, Guanyu Li, Zehui Chen, Zhengyin Du, Xuesong Yao, Yufei Xu, Jiecao Chen, Tao Gui, Zuxuan Wu, Qi Zhang, Xuanjing Huang, and Yungang Jiang. Agentgym-rl: Training LLM agents for long-horizon decision making through multi-turn reinforcement learning. *CoRR*, abs/2509.08755, 2025c. doi: 10.48550/ARXIV.2509.08755. URL <https://doi.org/10.48550/arXiv.2509.08755>.
- Can Xu, Qingfeng Sun, Kai Zheng, Xiubo Geng, Pu Zhao, Jiazhan Feng, Chongyang Tao, and Daxin Jiang. Wizardlm: Empowering large language models to follow complex instructions. *arXiv preprint arXiv:2304.12244*, 2023.
- Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. Webshop: Towards scalable real-world web interaction with grounded language agents. *Advances in Neural Information Processing Systems*, 35:20744–20757, 2022a.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik R Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. In *The eleventh international conference on learning representations*, 2022b.
- Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. Dapo: An open-source llm reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025.
- Daoguang Zan, Zhirong Huang, Wei Liu, Hanwu Chen, Linhao Zhang, Shulin Xin, Lu Chen, Qi Liu, Xiaojian Zhong, Aoyan Li, et al. Multi-swe-bench: A multilingual benchmark for issue resolving. *arXiv preprint arXiv:2504.02605*, 2025.

- Yirong Zeng, Xiao Ding, Yutai Hou, Yuxian Wang, Li Du, Juyi Dai, Qiuyang Ding, Duyu Tang, Dandan Tu, Weiwen Liu, Bing Qin, and Ting Liu. Tool zero: Training tool-augmented llms via pure RL from scratch. *CoRR*, abs/2511.01934, 2025a. doi: 10.48550/ARXIV.2511.01934. URL <https://doi.org/10.48550/arXiv.2511.01934>.
- Yongcheng Zeng, Xinyu Cui, Xuanfa Jin, Guoqing Liu, Zexu Sun, Dong Li, Ning Yang, Jianye Hao, Haifeng Zhang, and Jun Wang. Evolving llms’ self-refinement capability via iterative preference optimization. *arXiv preprint arXiv:2502.05605*, 2025b.
- Simon Zhai, Hao Bai, Zipeng Lin, Jiayi Pan, Peter Tong, Yifei Zhou, Alane Suhr, Saining Xie, Yann LeCun, Yi Ma, et al. Fine-tuning large vision-language models as decision-making agents via reinforcement learning. *Advances in neural information processing systems*, 37:110935–110971, 2024.
- Yuanzhao Zhai, Tingkai Yang, Kele Xu, Dawei Feng, Cheng Yang, Bo Ding, and Huaimin Wang. Enhancing decision-making for LLM agents via step-level q-value models. In Toby Walsh, Julie Shah, and Zico Kolter, editors, *AAAI-25, Sponsored by the Association for the Advancement of Artificial Intelligence, February 25 - March 4, 2025, Philadelphia, PA, USA*, pages 27161–27169. AAAI Press, 2025. doi: 10.1609/AAAI.V39I25.34924. URL <https://doi.org/10.1609/aaai.v39i25.34924>.
- Enci Zhang, Xingang Yan, Wei Lin, Tianxiang Zhang, and Lu Qianchun. Learning like humans: Advancing llm reasoning capabilities via adaptive difficulty curriculum learning and expert-guided self-reformulation. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 6630–6644, 2025a.
- Michael J. Q. Zhang, W. Bradley Knox, and Eunsol Choi. Modeling future conversation turns to teach llms to ask clarifying questions. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025*. OpenReview.net, 2025b. URL <https://openreview.net/forum?id=cwuSAR7EKd>.
- Zhihao Zhang, Qiaole Dong, Qi Zhang, Jun Zhao, Enyu Zhou, Zhiheng Xi, Senjie Jin, Xiaoran Fan, Yuhao Zhou, Mingqi Wu, et al. Why reinforcement fine-tuning enables mllms preserve prior knowledge better: A data perspective. *arXiv preprint arXiv:2506.23508*, 2025c.
- Haiyan Zhao, Hanjie Chen, Fan Yang, Ninghao Liu, Huiqi Deng, Hengyi Cai, Shuaiqiang Wang, Dawei Yin, and Mengnan Du. Explainability for large language models: A survey. *ACM Trans. Intell. Syst. Technol.*, 15(2):20:1–20:38, 2024. doi: 10.1145/3639372. URL <https://doi.org/10.1145/3639372>.
- Shuyan Zhou, Frank F Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, et al. Webarena: A realistic web environment for building autonomous agents. *arXiv preprint arXiv:2307.13854*, 2023.

Appendix

A. Limitations and Future Work

This paper adopts a new perspective to investigate how reinforcement fine-tuning affects the generalization ability of LLM agents. Our experiments focus on the Qwen2.5 family and yield extensive empirical results together with substantive insights. Despite this progress, important challenges remain. First, we rely on default training/evaluation protocols and hyperparameter choices, without exhaustive tuning; more careful optimization may bring new findings. Second, we adopt GRPO—the most commonly used RL algorithm for LLM agents at present—rather than other approaches such as REINFORCE++ (Hu, 2025) or DAPO (Yu et al., 2025); we leave these for future work. Additionally, due to computational constraints, our sequential multi-environment training study does not enumerate all possible environment orderings; instead, we evaluate a small set of representative sequence. Although these experiments already reveal important insights, a broader and more fine-grained investigation of ordering effects is left for future work and may further substantiate our findings.

B. Dataset Details

We select five representative agent environments, including WebShop (Yao et al., 2022a), SearchQA (Dunn et al., 2017), TextCraft (Sanghi et al., 2022), Alf World (Shridhar et al., 2020), and BabyAI (Chevalier-Boisvert et al., 2018). The characteristics of each environment are shown in Table 1 in Section 3.2. Here, we provide their detailed types and action spaces in Table 4.

Table 4 | Detailed action spaces for each environment.

Environment	Types	Action Spaces
WebShop	Web Navigation	search, click
SearchQA	Q&A Search	search, answer
TextCraft	Text-based Game	get, craft, inventory
Alf World	Household	go to, open, close, take from, put in/on, use, heat, cool, clean, slice, inventory, look, examine
BabyAI	Embodied	turn right, turn left, move forward, go to, go through, toggle and go through, toggle, pickup, drop, check available actions

Our data is sourced from AgentGym (Xi et al., 2025b). Following practice in previous work (Bengio et al., 2009; Mukherjee et al., 2023), we categorize the tasks $\mathcal{U}$ into *easy* and *hard* difficulty levels (denoted as $\mathcal{U}_{\text{easy}}$ and $\mathcal{U}_{\text{hard}}$) based on the avg@8 results of the Qwen2.5-7B-Instruct model, while ensuring a balanced distribution of data between the two difficulty levels. The same categorization is applied to the test set. Detailed statistics for each environment are provided in Table 5.

Table 5 | Detailed data statistics for each environment.

Environment	Training dataset			Testing dataset		
	easy	hard	all	easy	hard	all
WebShop	2104	1826	3930	84	116	200
SearchQA	1960	2040	4000	120	280	400
TextCraft	235	209	444	57	43	100
Alf World	1267	1153	2420	55	145	200
BabyAI	398	412	810	52	38	90

C. Detailed algorithm of GRPO.

Group Relative Policy Optimization (GRPO) is an efficient online reinforcement learning algorithm tailored for LLMs. It eliminates the need for a separate critic network typically required in PPO, thereby reducing computational overhead. Instead, GRPO estimates the baseline using *group relative advantages* to significantly reduce gradient variance.

Specifically, for each input query q (derived from u), GRPO samples a group of outputs $\{y_1, y_2, \dots, y_G\}$ from the old policy $\pi_{\theta_{old}}$ and obtains their corresponding rewards $\{R_1, R_2, \dots, R_G\}$. The advantage A_i for each output is calculated by normalizing the rewards within the group:

$$A_i = \frac{r_i - \text{mean}(\{R_1, \dots, R_G\})}{\text{std}(\{R_1, \dots, R_G\})} \quad (4)$$

Finally, GRPO updates the policy by maximizing the following surrogate objective, which incorporates PPO-style clipping (Schulman et al., 2017) and a KL divergence penalty to ensure training stability:

$$\mathcal{J}_{\text{GRPO}}(\theta) = \mathbb{E} \left[\frac{1}{G} \sum_{i=1}^G \left(\min(r_i(\theta)A_i, \text{clip}(r_i(\theta), 1 - \epsilon, 1 + \epsilon)A_i) - \beta \mathbb{D}_{\text{KL}} \right) \right], \quad (5)$$

where $r_i(\theta) = \frac{\pi_{\theta}(y_i|q)}{\pi_{\theta_{old}}(y_i|q)}$ denotes the probability ratio between the new and old policies, $\{y_i\}_{i=1}^G \sim \pi_{\theta_{old}}(q)$ represents the outputs y_i sampled from the old policy $\pi_{\theta_{old}}$ given the query q , ϵ is the clipping parameter, and β is the coefficient for the KL divergence term.

D. Detailed Results of Sequential Cross-Environment Training




Figure 7 | Training dynamics of forgetting and transfer in other sequential two-stage cross-environment training with Qwen2.5-7B-Instruct, where blue and red denote the upstream environment and the downstream environment, respectively.

Figure 3 in Section 6 illustrates the training dynamics for 8 two-stage sequential training configurations. Here, we present the dynamics for the remaining 12 pairs in Figure 7. Additionally, detailed final results are reported in Table 6.

Table 6 | Results of sequential training across different environments with Qwen2.5-7B-Instruct model.

UPSTREAM	DOWNSTREAM	WEBSHOP	SEARCHQA	TEXTCRAFT	ALFWORLD	BABYAI
BASE MODEL		28.59	31.19	33.63	26.56	67.00
WEBSHOP		86.50	33.28	40.75	24.13	79.21
	+ SEARCHQA	85.08	41.44	44.63	17.81	78.88
	+ TEXTCRAFT	86.32	36.81	82.50	26.63	81.68
	+ ALFWORLD	85.99	34.22	43.75	90.69	76.34
	+ BABYAI	86.38	33.97	49.25	25.50	87.74
SEARCHQA		47.07	46.12	35.25	16.75	80.33
	+ WEBSHOP	80.35	46.16	34.50	18.44	78.59
	+ TEXTCRAFT	49.31	46.19	78.63	26.06	78.03
	+ ALFWORLD	39.42	42.92	41.00	94.00	81.29
	+ BABYAI	44.36	45.44	25.00	14.06	89.62
TEXTCRAFT		38.30	32.19	80.88	31.50	77.95
	+ WEBSHOP	84.33	34.59	77.00	18.31	78.44
	+ SEARCHQA	39.17	43.16	71.63	19.69	78.97
	+ ALFWORLD	42.04	31.56	76.50	90.38	62.19
	+ BABYAI	23.00	33.37	74.88	22.81	87.55
ALFWORLD		34.31	29.59	36.13	92.00	72.91
	+ WEBSHOP	82.95	31.50	35.88	87.75	49.65
	+ SEARCHQA	53.04	44.66	43.50	90.63	77.86
	+ TEXTCRAFT	50.31	30.81	76.00	91.19	77.01
	+ BABYAI	25.68	28.31	39.25	85.56	75.59
BABYAI		10.25	29.41	39.25	28.13	88.79
	+ WEBSHOP	83.60	30.19	36.38	33.38	90.66
	+ SEARCHQA	33.71	38.84	40.63	24.06	89.37
	+ TEXTCRAFT	17.95	33.37	86.00	29.31	88.60
	+ ALFWORLD	13.60	29.88	41.25	92.50	89.32

E. Detailed Results of average turns and generated tokens in different environment

Figure 2 in section 4 reports the average number of interaction turns and generated tokens across different environments. The more detailed results are reported in Figure 8.

F. Failure Mode Analysis

Through detailed analysis, we summarize 8 common error types:

1. **Instruction Misinterpretation:** The agent fails to understand the prompt correctly, does not follow the instructions as indicated, or generates the wrong answer format.
2. **Guessing or Fabrication:** The agent relies too heavily on internal parameters or makes unsupported guesses, instead of utilizing environmental tools to gather necessary external information.
3. **Action Execution Failure:** The agent generates actions that cannot be interpreted by the environment or do not correspond to specific objects within the environment.




Figure 8 | Average generated tokens and average turn across different environments for Qwen2.5-3B-Instruct model trained with varying difficulties.

4. **Constraint Prioritization Failure:** The agent recognizes multiple constraints in the instructions but fails to evaluate their relative importance correctly, resulting in the violation of core constraints in favor of secondary ones.
5. **Confirmation Bias:** The agent becomes confident that it has found the correct answer or completed key steps, but does not proceed with further verification. Alternatively, it may forcefully apply other clues or validate erroneous information from the prompt.
6. **State or Memory Inconsistency:** The agent exhibits contradictions over time, forgetting tools it has recently invoked or results it has obtained. Completed steps may be repeated unnecessarily.
7. **Logic or Numerical Deficit:** The agent makes errors in logical or numerical reasoning.
8. **Termination Protocol Failure:** The agent fails to issue the environment-specific termination command or prematurely terminates before completing the task.

The results of both in-domain and out-of-domain evaluations are discussed in Section 7. Here, we present a detailed error classification in Figure 9.

G. Case Study

In this section, we present specific case studies. It is worth noting that while all interactions are inherently text-based, we provide visualizations for selected cases to enhance clarity. Additionally, due to the excessive number of interaction turns, we omit the majority of intermediate steps, highlighting only the pivotal moments in the figures.

Figure 10 illustrates a case where a model trained on SearchQA is evaluated on WebShop, comparing its generated trajectory against that of the base model on the same task. As observed in the case, the base model tends to blindly input the entire content of the instruction into the search bar, resulting in the retrieval of numerous irrelevant items. Furthermore, the base model exhibits incoherent decision-making under multiple constraints and struggles to accurately extract key information from the voluminous HTML content returned by the environment, leading to the selection of items from incorrect categories. In contrast, the model trained on SearchQA learns to formulate more flexible search queries, as well as perform efficient information extraction from complex results, thereby enabling it to successfully retrieve key information and select the correct item.

Figure 11 highlights a key reason why agents struggle to generalize to SearchQA, using model




Figure 9 | Failure mode distribution across all train-test environment combinations.

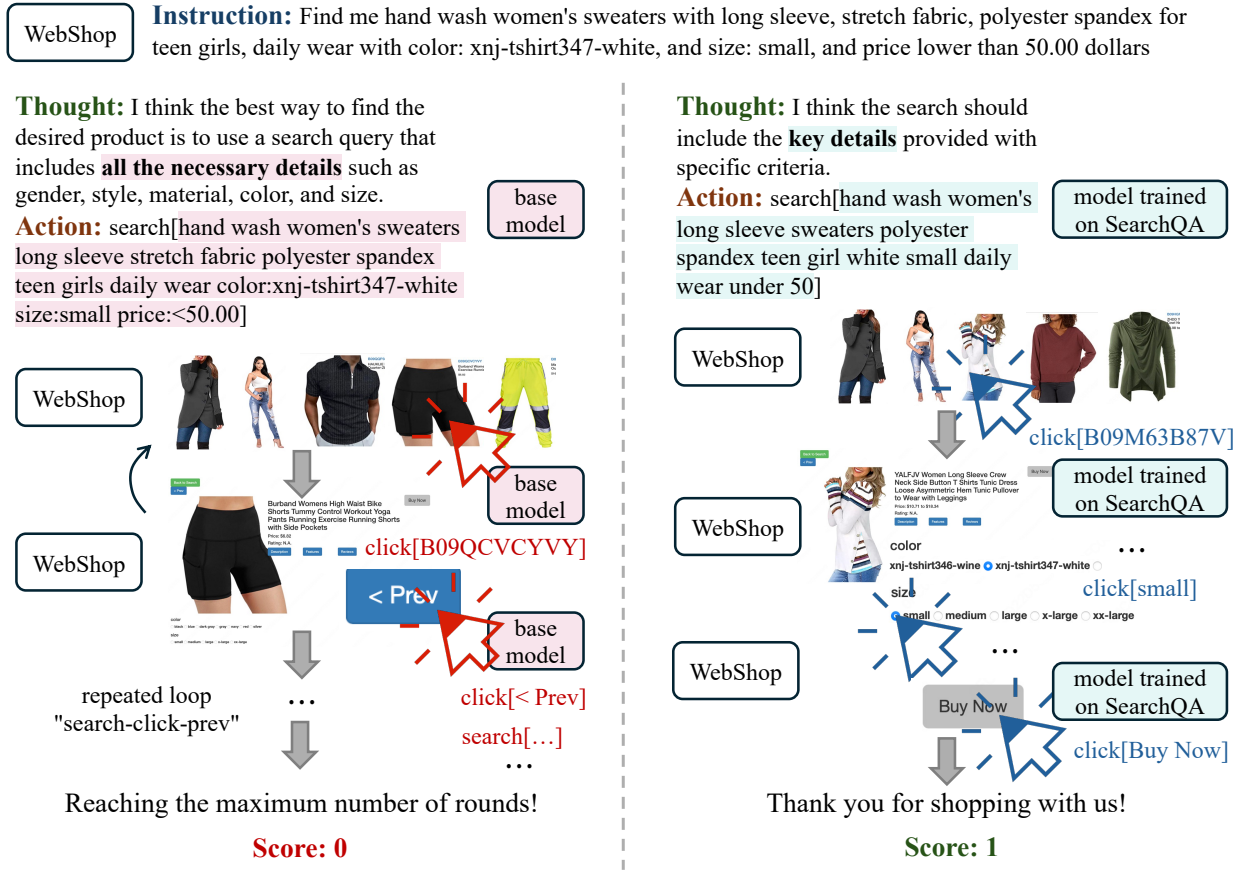


Figure 10 | Comparison of trajectories on WebShop between the base model and the model trained on SearchQA.

trained on AlfWorld as a case study. Although both models fail in their initial attempts, the model trained on SearchQA demonstrates the ability to refine its search queries to achieve greater precision. This capability proves difficult to transfer from other environments, e.g., the model trained on AlfWorld falls into a repetitive loop of the same searches and answers, consequently failing to resolve the task.

Figure 12 illustrates how training enhances exploration efficiency, using BabyAI as a case study. The base model, lacking strong spatial awareness, struggles to accurately pinpoint the target location based on textual descriptions. Consequently, it falls into a pattern of redundant exploration; although it eventually completes the task, it requires an excessive number of interaction turns. In contrast, the trained model is able to precisely locate the target, thereby completing the task via an optimal path and significantly improving exploration efficiency.

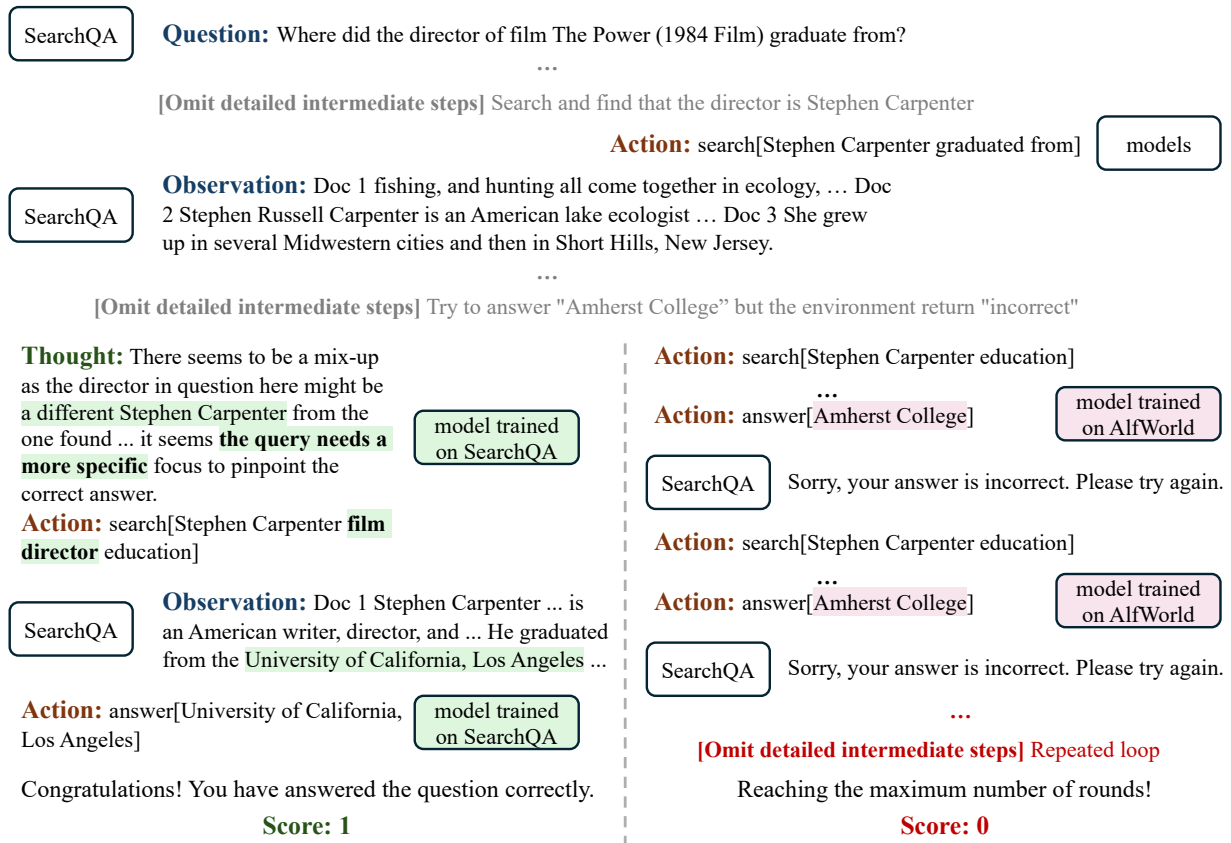


Figure 11 | Comparison of trajectories on SearchQA between the model trained on SearchQA and the mode trained on Alf World.

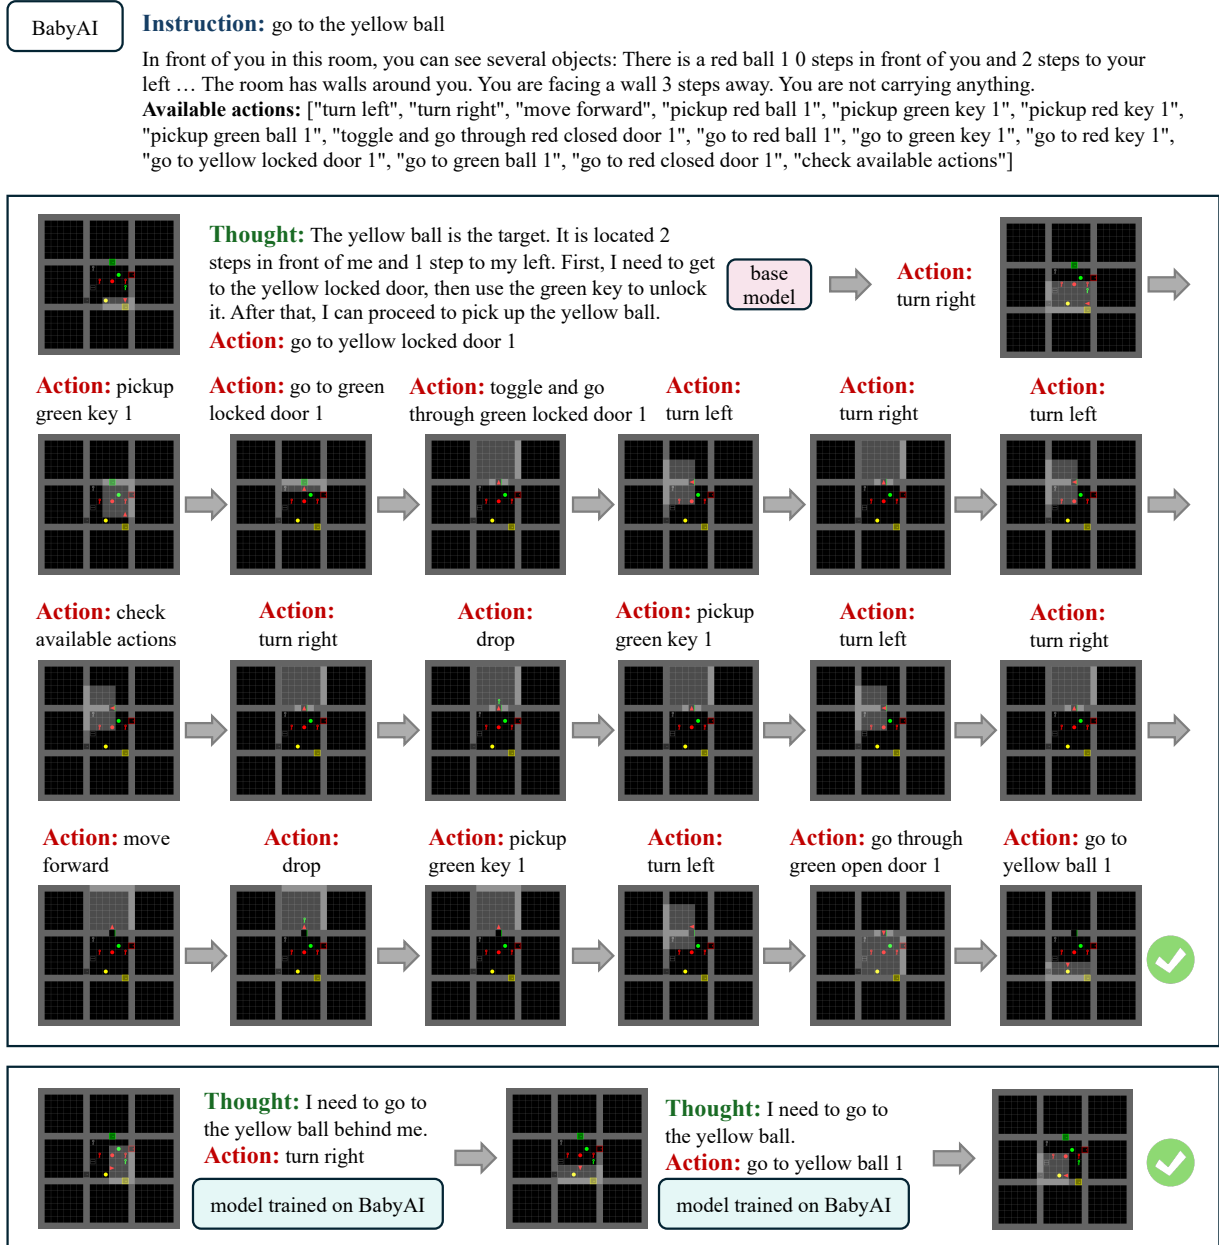


Figure 12 | A case study in the BabyAI environment. Compared to the base model, the trained model demonstrates significantly improved exploration efficiency.